

Mnacho Echenim

small evolutions : Patrick Reignier

Grenoble INP-Ensimag

2025-2026

Outline

More on backpropagation, mini-batch gradient descent

Improving gradient descent

Derivation of the backpropagation rules (simplified notations)

- ▶ **Goal:** get rid of cumbersome notations to represent partial derivatives

Derivation of the backpropagation rules (simplified notations)

- ▶ **Goal:** get rid of cumbersome notations to represent partial derivatives
- ▶ Closer to what may be found in textbooks

Derivation of the backpropagation rules (simplified notations)

- ▶ **Goal:** get rid of cumbersome notations to represent partial derivatives
- ▶ Closer to what may be found in textbooks
- ▶ We want to compute $\frac{\partial C}{\partial w_k^j}$ and $\frac{\partial C}{\partial b^j}$

Derivation of the backpropagation rules (simplified notations)

- ▶ **Goal:** get rid of cumbersome notations to represent partial derivatives
- ▶ Closer to what may be found in textbooks
- ▶ We want to compute $\frac{\partial C}{\partial w_k^j}$ and $\frac{\partial C}{\partial b^j}$
- ▶ We have the following equalities:

$$\frac{\partial C}{\partial w_k^j} = \frac{\partial C}{\partial a^L} \cdot \frac{\partial a^L}{\partial w_k^j}$$

Derivation of the backpropagation rules (simplified notations)

- ▶ **Goal:** get rid of cumbersome notations to represent partial derivatives
- ▶ Closer to what may be found in textbooks
- ▶ We want to compute $\frac{\partial C}{\partial w_k^j}$ and $\frac{\partial C}{\partial b^j}$
- ▶ We have the following equalities:

$$\begin{aligned}\frac{\partial C}{\partial w_k^j} &= \frac{\partial C}{\partial a^L} \cdot \frac{\partial a^L}{\partial w_k^j} \\ &= \frac{\partial C}{\partial a^L} \cdot \left[\prod_{i=L}^{j+1} \frac{\partial a^i}{\partial a^{i-1}} \right] \cdot \frac{\partial a^j}{\partial w_k^j}\end{aligned}$$

Derivation of the backpropagation rules (simplified notations)

- ▶ **Goal:** get rid of cumbersome notations to represent partial derivatives
- ▶ Closer to what may be found in textbooks
- ▶ We want to compute $\frac{\partial C}{\partial w_k^j}$ and $\frac{\partial C}{\partial b^j}$
- ▶ We have the following equalities:

$$\begin{aligned}
 \frac{\partial C}{\partial w_k^j} &= \frac{\partial C}{\partial a^L} \cdot \frac{\partial a^L}{\partial w_k^j} \\
 &= \frac{\partial C}{\partial a^L} \cdot \left[\prod_{i=L}^{j+1} \frac{\partial a^i}{\partial a^{i-1}} \right] \cdot \frac{\partial a^j}{\partial w_k^j} \\
 &= \mathcal{P}^{j+1} \cdot \frac{\partial a^j}{\partial w_k^j}
 \end{aligned}$$

Derivation of the backpropagation rules (simplified notations)

- ▶ **Goal:** get rid of cumbersome notations to represent partial derivatives
- ▶ Closer to what may be found in textbooks
- ▶ We want to compute $\frac{\partial \mathcal{C}}{\partial w_k^j}$ and $\frac{\partial \mathcal{C}}{\partial b^j}$
- ▶ We have the following equalities:

$$\begin{aligned}
 \frac{\partial \mathcal{C}}{\partial w_k^j} &= \frac{\partial \mathcal{C}}{\partial a^L} \cdot \frac{\partial a^L}{\partial w_k^j} \\
 &= \frac{\partial \mathcal{C}}{\partial a^L} \cdot \left[\prod_{i=L}^{j+1} \frac{\partial a^i}{\partial a^{i-1}} \right] \cdot \frac{\partial a^j}{\partial w_k^j} \\
 &= \mathcal{P}^{j+1} \cdot \frac{\partial a^j}{\partial w_k^j} \\
 \frac{\partial \mathcal{C}}{\partial b^j} &= \mathcal{P}^{j+1} \cdot \frac{\partial a^j}{\partial b^j}
 \end{aligned}$$

Derivation of the backpropagation rules (part 2)

$$\frac{\partial \mathbf{a}^j}{\partial \mathbf{a}^{j-1}} = \begin{pmatrix} \Phi'(\zeta_1^j) \cdot [\omega_1^j]^T \\ \vdots \\ \Phi'(\zeta_{n_j}^j) \cdot [\omega_{n_j}^j]^T \end{pmatrix}$$

Derivation of the backpropagation rules (part 2)

$$\frac{\partial \alpha^j}{\partial \alpha^{j-1}} = \begin{pmatrix} \Phi'(\zeta_1^j) \cdot [\omega_1^j]^T \\ \vdots \\ \Phi'(\zeta_{n_j}^j) \cdot [\omega_{n_j}^j]^T \end{pmatrix}$$

$$\frac{\partial \alpha^j}{\partial w_k^j} = \begin{pmatrix} 0 \\ \vdots \\ \Phi'(\zeta_k^j) \cdot [\alpha^{j-1}]^T \\ \vdots \\ 0 \end{pmatrix} \leftarrow \text{line } k$$

Derivation of the backpropagation rules (part 2)

$$\frac{\partial \mathbf{a}^j}{\partial \mathbf{a}^{j-1}} = \begin{pmatrix} \Phi'(\zeta_1^j) \cdot [\omega_1^j]^T \\ \vdots \\ \Phi'(\zeta_{n_j}^j) \cdot [\omega_{n_j}^j]^T \end{pmatrix}$$

$$\frac{\partial \mathbf{a}^j}{\partial w_k^j} = \begin{pmatrix} 0 \\ \vdots \\ \Phi'(\zeta_k^j) \cdot [\alpha^{j-1}]^T \\ \vdots \\ 0 \end{pmatrix} \leftarrow \text{line } k$$

$$\frac{\partial \mathbf{a}^j}{\partial b^j} = \text{diag} (\Phi'(\zeta_1^j), \dots, \Phi'(\zeta_{n_j}^j))$$

Derivation of the backpropagation rules (part 2)

$$\frac{\partial \mathbf{a}^j}{\partial \mathbf{a}^{j-1}} = \begin{pmatrix} \Phi'(\zeta_1^j) \cdot [\omega_1^j]^T \\ \vdots \\ \Phi'(\zeta_{n_j}^j) \cdot [\omega_{n_j}^j]^T \end{pmatrix}$$

$$\frac{\partial \mathbf{a}^j}{\partial w_k^j} = \begin{pmatrix} 0 \\ \vdots \\ \Phi'(\zeta_k^j) \cdot [\alpha^{j-1}]^T \\ \vdots \\ 0 \end{pmatrix} \leftarrow \text{line } k$$

$$\frac{\partial \mathbf{a}^j}{\partial b^j} = \text{diag} (\Phi'(\zeta_1^j), \dots, \Phi'(\zeta_{n_j}^j))$$

The backpropagation equations are then derived as previously

Mini-batch gradient descent

- ▶ We are given M inputs $\overline{\alpha^0} \stackrel{\text{def}}{=} (\alpha_{(1)}^0, \dots, \alpha_{(M)}^0)$ and outputs $\overline{\rho} \stackrel{\text{def}}{=} (\rho_{(1)}, \dots, \rho_{(M)})$

Mini-batch gradient descent

- ▶ We are given M inputs $\overline{\alpha^0} \stackrel{\text{def}}{=} (\alpha_{(1)}^0, \dots, \alpha_{(M)}^0)$ and outputs $\overline{\rho} \stackrel{\text{def}}{=} (\rho_{(1)}, \dots, \rho_{(M)})$
- ▶ The parameters updates become

$$\Omega^j \leftarrow \Omega^j - \frac{\eta}{M} \cdot \sum_{k=1}^M \nabla_{W^j} \mathcal{E}(\alpha_{(k)}^0, \dots, \Omega^L, \beta^L, \rho_{(k)})$$

Mini-batch gradient descent

- ▶ We are given M inputs $\overline{\alpha^0} \stackrel{\text{def}}{=} (\alpha_{(1)}^0, \dots, \alpha_{(M)}^0)$ and outputs $\overline{\rho} \stackrel{\text{def}}{=} (\rho_{(1)}, \dots, \rho_{(M)})$
- ▶ The parameters updates become

$$\Omega^j \leftarrow \Omega^j - \frac{\eta}{M} \cdot \sum_{k=1}^M \nabla_{W^j} \mathcal{E}(\alpha_{(k)}^0, \dots, \Omega^L, \beta^L, \rho_{(k)})$$

$$\Omega^j \leftarrow \Omega^j - \frac{\eta}{M} \cdot \sum_{k=1}^M \nabla_{W^j} \mathcal{C}(\alpha_{(k)}^L, \rho_{(k)})$$

Mini-batch gradient descent

- ▶ We are given M inputs $\overline{\alpha^0} \stackrel{\text{def}}{=} (\alpha_{(1)}^0, \dots, \alpha_{(M)}^0)$ and outputs $\overline{\rho} \stackrel{\text{def}}{=} (\rho_{(1)}, \dots, \rho_{(M)})$
- ▶ The parameters updates become

$$\Omega^j \leftarrow \Omega^j - \frac{\eta}{M} \cdot \sum_{k=1}^M \nabla_{W^j} \mathcal{E}(\alpha_{(k)}^0, \dots, \Omega^L, \beta^L, \rho_{(k)})$$

$$\Omega^j \leftarrow \Omega^j - \frac{\eta}{M} \cdot \sum_{k=1}^M \nabla_{W^j} \mathcal{C}(\alpha_{(k)}^L, \rho_{(k)})$$

$$\beta^j \leftarrow \beta^j - \frac{\eta}{M} \cdot \sum_{k=1}^M \nabla_{b^j} \mathcal{E}(\alpha_{(k)}^0, \dots, \Omega^L, \beta^L, \rho_{(k)})$$

Mini-batch gradient descent

- We are given M inputs $\overline{\alpha^0} \stackrel{\text{def}}{=} (\alpha_{(1)}^0, \dots, \alpha_{(M)}^0)$ and outputs $\overline{\rho} \stackrel{\text{def}}{=} (\rho_{(1)}, \dots, \rho_{(M)})$
- The parameters updates become

$$\Omega^j \leftarrow \Omega^j - \frac{\eta}{M} \cdot \sum_{k=1}^M \nabla_{W^j} \mathcal{E}(\alpha_{(k)}^0, \dots, \Omega^L, \beta^L, \rho_{(k)})$$

$$\Omega^j \leftarrow \Omega^j - \frac{\eta}{M} \cdot \sum_{k=1}^M \nabla_{W^j} \mathcal{C}(\alpha_{(k)}^L, \rho_{(k)})$$

$$\beta^j \leftarrow \beta^j - \frac{\eta}{M} \cdot \sum_{k=1}^M \nabla_{b^j} \mathcal{E}(\alpha_{(k)}^0, \dots, \Omega^L, \beta^L, \rho_{(k)})$$

$$\beta^j \leftarrow \beta^j - \frac{\eta}{M} \cdot \sum_{k=1}^M \nabla_{b^j} \mathcal{C}(\alpha_{(k)}^L, \rho_{(k)})$$

Mini-batch gradient descent

- ▶ We are given M inputs $\overline{\alpha^0} \stackrel{\text{def}}{=} (\alpha_{(1)}^0, \dots, \alpha_{(M)}^0)$ and outputs $\overline{\rho} \stackrel{\text{def}}{=} (\rho_{(1)}, \dots, \rho_{(M)})$
- ▶ The parameters updates become

$$\begin{aligned}\Omega^j &\leftarrow \Omega^j - \frac{\eta}{M} \cdot \sum_{k=1}^M \nabla_{W^j} \mathcal{E}(\alpha_{(k)}^0, \dots, \Omega^L, \beta^L, \rho_{(k)}) \\ \Omega^j &\leftarrow \Omega^j - \frac{\eta}{M} \cdot \sum_{k=1}^M \nabla_{W^j} \mathcal{C}(\alpha_{(k)}^L, \rho_{(k)}) \\ \beta^j &\leftarrow \beta^j - \frac{\eta}{M} \cdot \sum_{k=1}^M \nabla_{b^j} \mathcal{E}(\alpha_{(k)}^0, \dots, \Omega^L, \beta^L, \rho_{(k)}) \\ \beta^j &\leftarrow \beta^j - \frac{\eta}{M} \cdot \sum_{k=1}^M \nabla_{b^j} \mathcal{C}(\alpha_{(k)}^L, \rho_{(k)})\end{aligned}$$

- ▶ **Goal:** compute these updates efficiently using matrix operations

Extensions to mini-batches

Inputs to the network: $\overline{\alpha^0}, \Omega_1, \beta^1, \dots, \Omega^L, \beta^L$. We let:

$$\begin{aligned}
 \overline{\beta^i} &\stackrel{\text{def}}{=} (\beta^i, \dots, \beta^i) \quad (M \text{ columns}) \\
 \overline{\alpha^i} &\stackrel{\text{def}}{=} (\alpha_{(1)}^i, \dots, \alpha_{(M)}^i) = f_i(\overline{\alpha^{i-1}}, \Omega^i, \beta^i) \\
 \overline{\zeta^i} &\stackrel{\text{def}}{=} (\zeta_{(1)}^i, \dots, \zeta_{(M)}^i) \\
 \Phi'(\overline{\zeta^i}) &\stackrel{\text{def}}{=} (\Phi'(\zeta_{(1)}^i), \dots, \Phi'(\zeta_{(M)}^i)) \\
 \overline{\mathcal{B}^i} &\stackrel{\text{def}}{=} (\mathcal{B}_{(1)}^i, \dots, \mathcal{B}_{(M)}^i) \\
 \nabla_{\alpha^L} \mathcal{C}(\overline{\alpha^L}, \bar{\rho}) &= (\nabla_{\alpha_{(1)}^L} \mathcal{C}(\alpha_{(1)}^L, \rho_{(1)}), \dots, \nabla_{\alpha_{(M)}^L} \mathcal{C}(\alpha_{(M)}^L, \rho_{(M)}))
 \end{aligned}$$

Extensions to mini-batches

Inputs to the network: $\overline{\alpha^0}, \Omega_1, \beta^1, \dots, \Omega^L, \beta^L$. We let:

$$\begin{aligned}\overline{\beta^i} &\stackrel{\text{def}}{=} (\beta^i, \dots, \beta^i) \quad (M \text{ columns}) \\ \overline{\alpha^i} &\stackrel{\text{def}}{=} (\alpha_{(1)}^i, \dots, \alpha_{(M)}^i) = f_i(\overline{\alpha^{i-1}}, \Omega^i, \beta^i) \\ \overline{\zeta^i} &\stackrel{\text{def}}{=} (\zeta_{(1)}^i, \dots, \zeta_{(M)}^i) \\ \Phi'(\overline{\zeta^i}) &\stackrel{\text{def}}{=} (\Phi'(\zeta_{(1)}^i), \dots, \Phi'(\zeta_{(M)}^i)) \\ \overline{\mathcal{B}^i} &\stackrel{\text{def}}{=} (\mathcal{B}_{(1)}^i, \dots, \mathcal{B}_{(M)}^i) \\ \nabla_{\alpha^L} \mathcal{C}(\overline{\alpha^L}, \bar{\rho}) &= (\nabla_{\alpha_{(1)}^L} \mathcal{C}(\alpha_{(1)}^L, \rho_{(1)}), \dots, \nabla_{\alpha_{(M)}^L} \mathcal{C}(\alpha_{(M)}^L, \rho_{(M)}))\end{aligned}$$

We have the following equalities:

- ▶ $\overline{\zeta^i} = \Psi(\overline{\alpha^{i-1}}, \Omega^i, \overline{\beta^i}) = [\Omega^i]^T \cdot \overline{\alpha^{i-1}} + \overline{\beta^i}$ for $i = 1, \dots, L$
- ▶ $\overline{\mathcal{B}^L} = \Phi'(\overline{\zeta^L}) \odot \nabla_{\alpha^L} \mathcal{C}(\overline{\alpha^L}, \bar{\rho})$
- ▶ $\overline{\mathcal{B}^j} = \Phi'(\overline{\zeta^j}) \odot (\Omega^{j+1} \cdot \overline{\mathcal{B}^{j+1}})$

Other computations

$$\nabla_{W^j} \mathcal{C} = \frac{1}{M} \cdot \sum_{k=1}^M \nabla_{W^j} \mathcal{C}(\alpha_{(k)}^L, \rho_{(k)})$$

Other computations

$$\begin{aligned}\nabla_{W^j} \mathcal{C} &= \frac{1}{M} \cdot \sum_{k=1}^M \nabla_{W^j} \mathcal{C}(\alpha_{(k)}^L, \rho_{(k)}) \\ &= \frac{1}{M} \cdot \sum_{k=1}^M \alpha_{(k)}^{j-1} \cdot \left[\mathcal{B}_{(k)}^j \right]^T\end{aligned}$$

Other computations

$$\begin{aligned}\nabla_{W^j} \mathcal{C} &= \frac{1}{M} \cdot \sum_{k=1}^M \nabla_{W^j} \mathcal{C}(\alpha_{(k)}^L, \rho_{(k)}) \\ &= \frac{1}{M} \cdot \sum_{k=1}^M \alpha_{(k)}^{j-1} \cdot \left[\mathcal{B}_{(k)}^j \right]^T = \frac{1}{M} \cdot \overline{\alpha^{j-1}} \cdot \left[\overline{\mathcal{B}^j} \right]^T\end{aligned}$$

Other computations

$$\begin{aligned}\nabla_{W^j} \mathcal{C} &= \frac{1}{M} \cdot \sum_{k=1}^M \nabla_{W^j} \mathcal{C}(\alpha_{(k)}^L, \rho_{(k)}) \\ &= \frac{1}{M} \cdot \sum_{k=1}^M \alpha_{(k)}^{j-1} \cdot [\mathcal{B}_{(k)}^j]^T = \frac{1}{M} \cdot \overline{\alpha^{j-1}} \cdot [\overline{\mathcal{B}^j}]^T \\ \nabla_{b^j} \mathcal{C} &= \frac{1}{M} \cdot \sum_{k=1}^M \mathcal{B}_{(k)}^j\end{aligned}$$

Other computations

$$\begin{aligned}\nabla_{W^j} \mathcal{C} &= \frac{1}{M} \cdot \sum_{k=1}^M \nabla_{W^j} \mathcal{C}(\alpha_{(k)}^L, \rho_{(k)}) \\ &= \frac{1}{M} \cdot \sum_{k=1}^M \alpha_{(k)}^{j-1} \cdot [\mathcal{B}_{(k)}^j]^T = \frac{1}{M} \cdot \overline{\alpha^{j-1}} \cdot [\overline{\mathcal{B}^j}]^T \\ \nabla_{b^j} \mathcal{C} &= \frac{1}{M} \cdot \sum_{k=1}^M \mathcal{B}_{(k)}^j = \frac{1}{M} \cdot (\overline{\mathcal{B}^j} \cdot \mathbf{1}_M)\end{aligned}$$

Other computations

$$\begin{aligned}
 \nabla_{W^j} \mathcal{C} &= \frac{1}{M} \cdot \sum_{k=1}^M \nabla_{W^j} \mathcal{C}(\alpha_{(k)}^L, \rho_{(k)}) \\
 &= \frac{1}{M} \cdot \sum_{k=1}^M \alpha_{(k)}^{j-1} \cdot [\mathcal{B}_{(k)}^j]^T = \frac{1}{M} \cdot \overline{\alpha^{j-1}} \cdot [\overline{\mathcal{B}^j}]^T \\
 \nabla_{b^j} \mathcal{C} &= \frac{1}{M} \cdot \sum_{k=1}^M \mathcal{B}_{(k)}^j = \frac{1}{M} \cdot (\overline{\mathcal{B}^j} \cdot \mathbf{1}_M)
 \end{aligned}$$

Let $\mathbf{1}_M \stackrel{\text{def}}{=} (1, \dots, 1)^T \in \mathbb{R}^M$. For M inputs and outputs, parameter updates become:

$$\begin{aligned}
 \Omega^j &\leftarrow \Omega^j - \frac{\eta}{M} \cdot \left(\overline{\alpha^{j-1}} \cdot [\overline{\mathcal{B}^j}]^T \right) \\
 \beta_j &\leftarrow \beta_j - \frac{\eta}{M} \cdot (\overline{\mathcal{B}^j} \cdot \mathbf{1}_M)
 \end{aligned}$$

Mini-batch forward and backpropagation algorithms

Input: A network with L layers

Input: $\overline{\alpha^0} = (\alpha_{(1)}^0, \dots, \alpha_{(M)}^0)$

1 **for** $i \leftarrow 1$ **to** L **do**

2 $\overline{\zeta^i} \leftarrow [\Omega^i]^T \cdot \overline{\alpha^{i-1}} + \overline{\beta^i};$

3 $\overline{\alpha^i} \leftarrow \Phi(\overline{\zeta^i});$

4 **end**

Algorithm 1: Forward propagation

Mini-batch forward and backpropagation algorithms

Input: A network with L layers

Input: $\overline{\alpha^0} = (\alpha_{(1)}^0, \dots, \alpha_{(M)}^0)$

```

1 for  $i \leftarrow 1$  to  $L$  do
2    $\overline{\zeta^i} \leftarrow [\Omega^i]^T \cdot \overline{\alpha^{i-1}} + \overline{\beta^i};$ 
3    $\overline{\alpha^i} \leftarrow \Phi(\overline{\zeta^i});$ 
4 end
```

Algorithm 3: Forward propagation

Input: A network with L layers

Input: $\overline{\alpha^0} = (\alpha_{(1)}^0, \dots, \alpha_{(M)}^0)$ that has been forward propagated

Input: $\overline{\rho} = (\rho_{(1)}, \dots, \rho_{(M)})$ the expected outputs

```

1  $\overline{\mathcal{B}^L} \leftarrow \Phi'(\overline{\zeta^L}) \odot \nabla_{\alpha^L} \mathcal{C}(\overline{\alpha^L}, \overline{\rho});$ 
2  $[\overline{\mathcal{P}^L}]^T \leftarrow \Omega^L \cdot \overline{\mathcal{B}^L};$ 
3 for  $j \leftarrow L - 1$  to 1 do
4    $\overline{\mathcal{B}^j} \leftarrow \Phi'(\overline{\zeta^j}) \odot [\overline{\mathcal{P}^{j+1}}]^T;$ 
5    $[\overline{\mathcal{P}^j}]^T \leftarrow \Omega^j \cdot \overline{\mathcal{B}^j};$ 
6 end
```

Algorithm 4: Backpropagation

Mini-batch gradient computation

Input: A network with L layers

Input: $\overline{\alpha^0} = (\alpha_{(1)}^0, \dots, \alpha_{(M)}^0)$ that has been forward propagated

Input: $(\overline{\mathcal{B}^1}, \dots, \overline{\mathcal{B}^L})$ that have been updated by backpropagation

```
1 for  $j \in \{1, \dots, L\}$  do
2   Gradient( $\Omega^j$ )  $\leftarrow \frac{1}{M} \cdot \left( \overline{\alpha^{j-1}} \cdot [\overline{\mathcal{B}^j}]^T \right);$ 
3   Gradient( $\beta^j$ )  $\leftarrow \frac{1}{M} \cdot \left( \overline{\mathcal{B}^j} \cdot \mathbf{1}_M \right);$ 
4 end
```

Algorithm 5: Mini-batch gradient computation

Mini-batch gradient computation

Input: A network with L layers

Input: $\overline{\alpha^0} = (\alpha_{(1)}^0, \dots, \alpha_{(M)}^0)$ that has been forward propagated

Input: $(\overline{\mathcal{B}^1}, \dots, \overline{\mathcal{B}^L})$ that have been updated by backpropagation

```
1 for  $j \in \{1, \dots, L\}$  do
2   Gradient( $\Omega^j$ )  $\leftarrow \frac{1}{M} \cdot \left( \overline{\alpha^{j-1}} \cdot [\overline{\mathcal{B}^j}]^T \right);$ 
3   Gradient( $\beta^j$ )  $\leftarrow \frac{1}{M} \cdot \left( \overline{\mathcal{B}^j} \cdot \mathbf{1}_M \right);$ 
4 end
```

Algorithm 6: Mini-batch gradient computation

Note

By storing the necessary information, backpropagation and mini-batch gradient computations can be carried out in the same loop

Computed quantities in backpropagation

- ▶ At layer $j \in \llbracket 1, L \rrbracket$, two central quantities are computed
 - ▶ $\mathcal{P}^j \in \mathbb{R}^{1 \times n_{j-1}}$ and $\mathcal{B}^j \in \mathbb{R}^{n_j}$

Computed quantities in backpropagation

- ▶ At layer $j \in \llbracket 1, L \rrbracket$, two central quantities are computed
 - ▶ $\mathcal{P}^j \in \mathbb{R}^{1 \times n_{j-1}}$ and $\mathcal{B}^j \in \mathbb{R}^{n_j}$
 - ▶ \mathcal{P}^j represents the information that is transmitted to layer $j - 1$

Computed quantities in backpropagation

- ▶ At layer $j \in \llbracket 1, L \rrbracket$, two central quantities are computed
 - ▶ $\mathcal{P}^j \in \mathbb{R}^{1 \times n_{j-1}}$ and $\mathcal{B}^j \in \mathbb{R}^{n_j}$
- ▶ \mathcal{P}^j represents the information that is transmitted to layer $j - 1$
- ▶ \mathcal{B}^j represents the information necessary to compute gradients at layer j

Computed quantities in backpropagation

- ▶ At layer $j \in \llbracket 1, L \rrbracket$, two central quantities are computed
 - ▶ $\mathcal{P}^j \in \mathbb{R}^{1 \times n_{j-1}}$ and $\mathcal{B}^j \in \mathbb{R}^{n_j}$
 - ▶ \mathcal{P}^j represents the information that is transmitted to layer $j - 1$
 - ▶ \mathcal{B}^j represents the information necessary to compute gradients at layer j

- ▶ We have $\mathcal{P}^j = \frac{\partial \mathcal{C}}{\partial a^{j-1}}$
- ▶ If z^j is a formal parameter representing the net input of layer j , then
$$\mathcal{B}^j = \nabla_{z^j} \mathcal{C}$$

Outline

More on backpropagation, mini-batch gradient descent

Improving gradient descent

Foreword

- ▶ What follows is a list of techniques that are often used to improve gradient descent

Foreword

- ▶ What follows is a list of techniques that are often used to improve gradient descent
- ▶ There are (currently) no formal proofs that these methods improve anything on neural networks
 - ▶ They may be viewed by some as nothing more than recipes

Foreword

- ▶ What follows is a list of techniques that are often used to improve gradient descent
- ▶ There are (currently) no formal proofs that these methods improve anything on neural networks
 - ▶ They may be viewed by some as nothing more than recipes
- ▶ **But**
 - ▶ Several originate from research on improving gradient descent on convex optimization problems
 - ▶ E.g., the Nesterov momentum method permits to obtain optimal convergence rates for convex optimization problems
 - ▶ It is not far-fetched to try them on non-convex optimization problems

Reducing noise for mini-batch gradient descent

Issues with mini-batch gradient descent

- ▶ As we get closer to an optimum, the noise in the gradient estimate can become a problem
- ▶ Demo 1

Reducing noise for mini-batch gradient descent

Issues with mini-batch gradient descent

- ▶ As we get closer to an optimum, the noise in the gradient estimate can become a problem
- ▶ Demo 1

How can this issue be handled?

Reducing noise for mini-batch gradient descent

Issues with mini-batch gradient descent

- ▶ As we get closer to an optimum, the noise in the gradient estimate can become a problem
- ▶ Demo 1

How can this issue be handled?

- ▶ First idea: “denoise” the estimate by increasing the batch sizes

Reducing noise for mini-batch gradient descent

Issues with mini-batch gradient descent

- ▶ As we get closer to an optimum, the noise in the gradient estimate can become a problem
- ▶ Demo 1

How can this issue be handled?

- ▶ First idea: “denoise” the estimate by increasing the batch sizes
- ▶ Another idea: decrease the learning rate over time with decrease factor δ

Reducing noise for mini-batch gradient descent

Issues with mini-batch gradient descent

- ▶ As we get closer to an optimum, the noise in the gradient estimate can become a problem
- ▶ Demo 1

How can this issue be handled?

- ▶ First idea: “denoise” the estimate by increasing the batch sizes
- ▶ Another idea: decrease the learning rate over time with decrease factor δ
 - ▶ Inverse decay: $\eta_k = \frac{\eta_0}{1+k\delta}$

Reducing noise for mini-batch gradient descent

Issues with mini-batch gradient descent

- ▶ As we get closer to an optimum, the noise in the gradient estimate can become a problem
- ▶ Demo 1

How can this issue be handled?

- ▶ First idea: “denoise” the estimate by increasing the batch sizes
- ▶ Another idea: decrease the learning rate over time with decrease factor δ
 - ▶ Inverse decay: $\eta_k = \frac{\eta_0}{1+k\delta}$
 - ▶ Exponential decay: $\eta_k = \eta_0 \exp(-k\delta)$

Reducing noise for mini-batch gradient descent

Issues with mini-batch gradient descent

- ▶ As we get closer to an optimum, the noise in the gradient estimate can become a problem
- ▶ Demo 1

How can this issue be handled?

- ▶ First idea: “denoise” the estimate by increasing the batch sizes
- ▶ Another idea: decrease the learning rate over time with decrease factor δ
 - ▶ Inverse decay: $\eta_k = \frac{\eta_0}{1+k\delta}$
 - ▶ Exponential decay: $\eta_k = \eta_0 \exp(-k\delta)$
 - ▶ ...

Reducing noise for mini-batch gradient descent

Issues with mini-batch gradient descent

- ▶ As we get closer to an optimum, the noise in the gradient estimate can become a problem
- ▶ Demo 1

How can this issue be handled?

- ▶ First idea: “denoise” the estimate by increasing the batch sizes
- ▶ Another idea: decrease the learning rate over time with decrease factor δ
 - ▶ Inverse decay: $\eta_k = \frac{\eta_0}{1+k\delta}$
 - ▶ Exponential decay: $\eta_k = \eta_0 \exp(-k\delta)$
 - ▶ ...
- ▶ Demo 2

Reducing noise for mini-batch gradient descent

Issues with mini-batch gradient descent

- ▶ As we get closer to an optimum, the noise in the gradient estimate can become a problem
- ▶ Demo 1

How can this issue be handled?

- ▶ First idea: “denoise” the estimate by increasing the batch sizes
- ▶ Another idea: decrease the learning rate over time with decrease factor δ
 - ▶ Inverse decay: $\eta_k = \frac{\eta_0}{1+k\delta}$
 - ▶ Exponential decay: $\eta_k = \eta_0 \exp(-k\delta)$
 - ▶ ...
- ▶ Demo 2

Note

It is still difficult to choose the initial rate, decrease factor, decrease schedule...

Is it guaranteed gradient descent will work?

- ▶ Problem: are we sure gradient descent will lead to a global minimum of the function?

Is it guaranteed gradient descent will work?

- ▶ Problem: are we sure gradient descent will lead to a global minimum of the function?
- ▶ In general no: most of the time, the function to optimize is not convex

Is it guaranteed gradient descent will work?

- ▶ Problem: are we sure gradient descent will lead to a global minimum of the function?
- ▶ In general no: most of the time, the function to optimize is not convex
- ▶ Gradient descent could get stuck on
 - ▶ Local minima
 - ▶ Stationary points

Is it guaranteed gradient descent will work?

- ▶ Problem: are we sure gradient descent will lead to a global minimum of the function?
- ▶ In general no: most of the time, the function to optimize is not convex
- ▶ Gradient descent could get stuck on
 - ▶ Local minima
 - ▶ Stationary points
- ▶ Can the algorithm be adapted to produce better results?

Adaptive learning rates

- ▶ Problem: fixing the right learning rate is difficult, even using decay strategies

Adaptive learning rates

- ▶ Problem: fixing the right learning rate is difficult, even using decay strategies
- ▶ It is not clear whether the same learning rate should be used on every component

Adaptive learning rates

- ▶ Problem: fixing the right learning rate is difficult, even using decay strategies
- ▶ It is not clear whether the same learning rate should be used on every component
- ▶ Principle: use information about computed gradients to set learning rates individually for each component

Adaptive learning rates

- ▶ Problem: fixing the right learning rate is difficult, even using decay strategies
- ▶ It is not clear whether the same learning rate should be used on every component
- ▶ Principle: use information about computed gradients to set learning rates individually for each component
- ▶ First approach: delta-bar-delta (Jacobs, 88)

Adaptive learning rates

- ▶ Problem: fixing the right learning rate is difficult, even using decay strategies
- ▶ It is not clear whether the same learning rate should be used on every component
- ▶ Principle: use information about computed gradients to set learning rates individually for each component
- ▶ First approach: delta-bar-delta (Jacobs, 88)
 - ▶ Principle: check the sign of each partial derivative

Adaptive learning rates

- ▶ Problem: fixing the right learning rate is difficult, even using decay strategies
- ▶ It is not clear whether the same learning rate should be used on every component
- ▶ Principle: use information about computed gradients to set learning rates individually for each component
- ▶ First approach: delta-bar-delta (Jacobs, 88)
 - ▶ Principle: check the sign of each partial derivative
 - ▶ If it stays the same, then the optimization direction is correct: **increase the learning rate**

Adaptive learning rates

- ▶ Problem: fixing the right learning rate is difficult, even using decay strategies
- ▶ It is not clear whether the same learning rate should be used on every component
- ▶ Principle: use information about computed gradients to set learning rates individually for each component
- ▶ First approach: delta-bar-delta (Jacobs, 88)
 - ▶ Principle: check the sign of each partial derivative
 - ▶ If it stays the same, then the optimization direction is correct: **increase the learning rate**
 - ▶ Otherwise **decrease the learning rate**

Adaptive learning rates

- ▶ Problem: fixing the right learning rate is difficult, even using decay strategies
- ▶ It is not clear whether the same learning rate should be used on every component
- ▶ Principle: use information about computed gradients to set learning rates individually for each component
- ▶ First approach: delta-bar-delta (Jacobs, 88)
 - ▶ Principle: check the sign of each partial derivative
 - ▶ If it stays the same, then the optimization direction is correct: **increase the learning rate**
 - ▶ Otherwise **decrease the learning rate**
 - ▶ **Nb:** This approach can only be applied to full gradient descent
 - ▶ The noise in SGD can produce wrong signals

Adaptive learning rates

- ▶ Problem: fixing the right learning rate is difficult, even using decay strategies
- ▶ It is not clear whether the same learning rate should be used on every component
- ▶ Principle: use information about computed gradients to set learning rates individually for each component
- ▶ First approach: delta-bar-delta (Jacobs, 88)
 - ▶ Principle: check the sign of each partial derivative
 - ▶ If it stays the same, then the optimization direction is correct: **increase the learning rate**
 - ▶ Otherwise **decrease the learning rate**
 - ▶ **Nb:** This approach can only be applied to full gradient descent
 - ▶ The noise in SGD can produce wrong signals
- ▶ Upcoming approaches: all operations on vectors and matrices are **componentwise**

AdaGrad (Duchi et al, 2011)

- ▶ Principle: keep track of aggregated squared magnitude of gradients; use this to scale the individual learning rates
 - ▶ Large partial derivatives lead to a fast decrease, and vice-versa

AdaGrad (Duchi et al, 2011)

- ▶ Principle: keep track of aggregated squared magnitude of gradients; use this to scale the individual learning rates
 - ▶ Large partial derivatives lead to a fast decrease, and vice-versa
- ▶ This technique has good properties in a convex setting
 - ▶ But the fact that **all** gradients are accumulated may lead to an excessive decrease of the learning rate

AdaGrad (Duchi et al, 2011)

- ▶ Principle: keep track of aggregated squared magnitude of gradients; use this to scale the individual learning rates
 - ▶ Large partial derivatives lead to a fast decrease, and vice-versa
- ▶ This technique has good properties in a convex setting
 - ▶ But the fact that **all** gradients are accumulated may lead to an excessive decrease of the learning rate

Input: Step size η , numerical stabilizer δ

Input: Initial parameters θ

```
1 set gradient accumulation variable  $r$  to 0;  
2 for  $i \leftarrow 1$  to number of training steps do  
3    $g \leftarrow$  computation of  $\nabla_{\theta} \mathcal{C}(\theta)$  //full gradient, mini-batch...;  
4    $r \leftarrow r + (g \odot g);$   
5    $v \leftarrow -\frac{\eta}{\delta+\sqrt{r}} \odot g;$   
6    $\theta \leftarrow \theta + v;$   
7 end
```

AdaGrad (Duchi et al, 2011)

- ▶ Principle: keep track of aggregated squared magnitude of gradients; use this to scale the individual learning rates
 - ▶ Large partial derivatives lead to a fast decrease, and vice-versa
- ▶ This technique has good properties in a convex setting
 - ▶ But the fact that **all** gradients are accumulated may lead to an excessive decrease of the learning rate

Input: Step size η , numerical stabilizer δ

Input: Initial parameters θ

```

1 set gradient accumulation variable  $r$  to 0;
2 for  $i \leftarrow 1$  to number of training steps do
3    $g \leftarrow$  computation of  $\nabla_{\theta} \mathcal{C}(\theta)$  //full gradient, mini-batch...;
4    $r \leftarrow r + (g \odot g);$ 
5    $v \leftarrow -\frac{\eta}{\delta+\sqrt{r}} \odot g;$ 
6    $\theta \leftarrow \theta + v;$ 
7 end

```

- ▶ Demo 3

RMSProp (Hinton, 2012)

- ▶ Principle: accumulate squared magnitude of gradients, with an exponentially weighted moving average
 - ▶ Ancient history is discarded, as old gradients decay exponentially with time

RMSProp (Hinton, 2012)

- ▶ Principle: accumulate squared magnitude of gradients, with an exponentially weighted moving average
 - ▶ Ancient history is discarded, as old gradients decay exponentially with time
- ▶ Example: on a slope with a locally convex bowl:
 - ▶ AdaGrad will likely get stuck
 - ▶ RMSProp will forget about the past and get out of the bowl

RMSProp (Hinton, 2012)

- ▶ Principle: accumulate squared magnitude of gradients, with an exponentially weighted moving average
 - ▶ Ancient history is discarded, as old gradients decay exponentially with time
- ▶ Example: on a slope with a locally convex bowl:
 - ▶ AdaGrad will likely get stuck
 - ▶ RMSProp will forget about the past and get out of the bowl

Input: Step size η , numerical stabilizer δ

Input: Exponential decay rate $\rho \in [0, 1[$

Input: Initial parameters θ

```

1 set gradient accumulation variable  $r$  to 0;
2 for  $i \leftarrow 1$  to number of training steps do
3    $g \leftarrow$  computation of  $\nabla_{\theta}\mathcal{C}(\theta)$  //full gradient, mini-batch...;
4    $r \leftarrow \rho \cdot r + (1 - \rho) \cdot (g \odot g);$ 
5    $v \leftarrow -\frac{\eta}{\sqrt{\delta+r}} \odot g;$ 
6    $\theta \leftarrow \theta + v;$ 
7 end
  
```

RMSProp (Hinton, 2012)

- ▶ Principle: accumulate squared magnitude of gradients, with an exponentially weighted moving average
 - ▶ Ancient history is discarded, as old gradients decay exponentially with time
- ▶ Example: on a slope with a locally convex bowl:
 - ▶ AdaGrad will likely get stuck
 - ▶ RMSProp will forget about the past and get out of the bowl

Input: Step size η , numerical stabilizer δ

Input: Exponential decay rate $\rho \in [0, 1[$

Input: Initial parameters θ

```

1 set gradient accumulation variable  $r$  to 0;
2 for  $i \leftarrow 1$  to number of training steps do
3    $g \leftarrow$  computation of  $\nabla_{\theta}\mathcal{C}(\theta)$  //full gradient, mini-batch...;
4    $r \leftarrow \rho \cdot r + (1 - \rho) \cdot (g \odot g);$ 
5    $v \leftarrow -\frac{\eta}{\sqrt{\delta+r}} \odot g;$ 
6    $\theta \leftarrow \theta + v;$ 
7 end
  
```

- ▶ Demo 4

AdaDelta (Zeiler, 2012)

- ▶ Principle: moving average, similarly to RMSProp

AdaDelta (Zeiler, 2012)

- ▶ Principle: moving average, similarly to RMSProp
 - ▶ Main difference: **no input step size**
 - ▶ Step size is updated depending on previous changes

AdaDelta (Zeiler, 2012)

- ▶ Principle: moving average, similarly to RMSProp
 - ▶ Main difference: **no input step size**
 - ▶ Step size is updated depending on previous changes

Input: Numerical stabilizer δ

Input: Exponential decay rate $\rho \in [0, 1[$

Input: Initial parameters θ

```

1 set gradient accumulation variable  $r$  to 0;
2 set parameter update accumulation variable  $s$  to 0;
3 for  $i \leftarrow 1$  to number of training steps do
4    $g \leftarrow$  computation of  $\nabla_{\theta}\mathcal{C}(\theta)$  //full gradient, mini-batch...;
5    $r \leftarrow \rho \cdot r + (1 - \rho) \cdot (g \odot g);$ 
6    $v \leftarrow -\frac{\sqrt{\delta+s}}{\sqrt{\delta+r}} \odot g;$ 
7    $s \leftarrow \rho \cdot s + (1 - \rho) \cdot (v \odot v);$ 
8    $\theta \leftarrow \theta + v;$ 
9 end
```

AdaDelta (Zeiler, 2012)

- ▶ Principle: moving average, similarly to RMSProp
 - ▶ Main difference: **no input step size**
 - ▶ Step size is updated depending on previous changes

Input: Numerical stabilizer δ

Input: Exponential decay rate $\rho \in [0, 1[$

Input: Initial parameters θ

```

1 set gradient accumulation variable  $r$  to 0;
2 set parameter update accumulation variable  $s$  to 0;
3 for  $i \leftarrow 1$  to number of training steps do
4    $g \leftarrow$  computation of  $\nabla_{\theta}\mathcal{C}(\theta)$  //full gradient, mini-batch...;
5    $r \leftarrow \rho \cdot r + (1 - \rho) \cdot (g \odot g);$ 
6    $v \leftarrow -\frac{\sqrt{\delta+s}}{\sqrt{\delta+r}} \odot g;$ 
7    $s \leftarrow \rho \cdot s + (1 - \rho) \cdot (v \odot v);$ 
8    $\theta \leftarrow \theta + v;$ 
9 end

```

- ▶ Demo 5

Adam (Kingma & Ba, 2014)

- ▶ Can be viewed as a combination of RMSProp and Momentum (coming up later)
 - ▶ Momentum is added by considering an exponentially weighted moving average of gradients

Adam (Kingma & Ba, 2014)

- ▶ Can be viewed as a combination of RMSProp and Momentum (coming up later)
 - ▶ Momentum is added by considering an exponentially weighted moving average of gradients
- ▶ Main difference: a bias correction term is applied to the first and second moment variables

Adam (Kingma & Ba, 2014)

- ▶ Can be viewed as a combination of RMSProp and Momentum (coming up later)
 - ▶ Momentum is added by considering an exponentially weighted moving average of gradients
- ▶ Main difference: a bias correction term is applied to the first and second moment variables

Input: Step size η , numerical stabilizer δ

Input: Exponential decay rates $\rho_1, \rho_2 \in [0, 1[$

Input: Initial parameters θ

```

1 set first, second moment variable  $s, r$  to 0;
2 for  $i \leftarrow 1$  to number of training steps do
3    $g \leftarrow$  computation of  $\nabla_{\theta} \mathcal{C}(\theta)$  //full gradient, mini-batch...;
4    $s \leftarrow \rho_1 \cdot s + (1 - \rho_1) \cdot g;$ 
5    $r \leftarrow \rho_2 \cdot r + (1 - \rho_2) \cdot (g \odot g);$ 
6    $s' \leftarrow \frac{s}{1 - \rho_1^i};$ 
7    $r' \leftarrow \frac{r}{1 - \rho_2^i};$ 
8    $v \leftarrow -\eta \cdot \frac{s'}{\delta + \sqrt{r'}};$ 
9    $\theta \leftarrow \theta + v;$ 
10 end
```

Adam (Kingma & Ba, 2014)

- ▶ Can be viewed as a combination of RMSProp and Momentum (coming up later)
 - ▶ Momentum is added by considering an exponentially weighted moving average of gradients
- ▶ Main difference: a bias correction term is applied to the first and second moment variables

Input: Step size η , numerical stabilizer δ

Input: Exponential decay rates $\rho_1, \rho_2 \in [0, 1[$

Input: Initial parameters θ

```

1 set first, second moment variable  $s, r$  to 0;
2 for  $i \leftarrow 1$  to number of training steps do
3    $g \leftarrow$  computation of  $\nabla_{\theta} \mathcal{C}(\theta)$  //full gradient, mini-batch...;
4    $s \leftarrow \rho_1 \cdot s + (1 - \rho_1) \cdot g;$ 
5    $r \leftarrow \rho_2 \cdot r + (1 - \rho_2) \cdot (g \odot g);$ 
6    $s' \leftarrow \frac{s}{1 - \rho_1^i};$ 
7    $r' \leftarrow \frac{r}{1 - \rho_2^i};$ 
8    $v \leftarrow -\eta \cdot \frac{s'}{\delta + \sqrt{r'}};$ 
9    $\theta \leftarrow \theta + v;$ 
10 end
```

- ▶ Demo 6

Momentum

- ▶ Principle: consider a *velocity* variable v

Momentum

- ▶ Principle: consider a *velocity* variable v
 - ▶ **Intuition:** v provides the direction and speed at which θ should move toward the optimum

Momentum

- ▶ Principle: consider a *velocity* variable v
 - ▶ **Intuition:** v provides the direction and speed at which θ should move toward the optimum
 - ▶ **How:** by using information from previous gradient computations

Momentum

- ▶ Principle: consider a *velocity* variable v
 - ▶ **Intuition:** v provides the direction and speed at which θ should move toward the optimum
 - ▶ **How:** by using information from previous gradient computations
 - ▶ For $\delta \in [0, 1[$, at each epoch, v is updated by $v \leftarrow \delta \cdot v - \eta \nabla_{\theta} \mathcal{C}(\theta)$

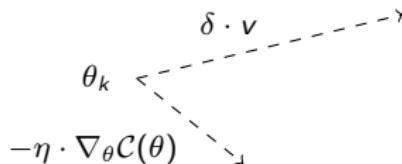
Momentum

- ▶ Principle: consider a *velocity* variable v
 - ▶ **Intuition:** v provides the direction and speed at which θ should move toward the optimum
 - ▶ **How:** by using information from previous gradient computations
 - ▶ For $\delta \in [0, 1[$, at each epoch, v is updated by $v \leftarrow \delta \cdot v - \eta \nabla_{\theta} \mathcal{C}(\theta)$

$$\begin{array}{c} \theta_k \\ \searrow \\ -\eta \cdot \nabla_{\theta} \mathcal{C}(\theta) \end{array}$$

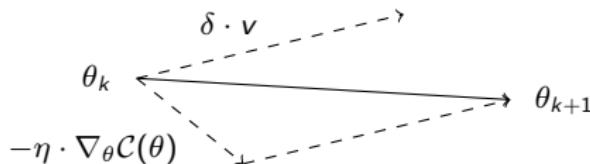
Momentum

- ▶ Principle: consider a *velocity* variable v
 - ▶ **Intuition:** v provides the direction and speed at which θ should move toward the optimum
 - ▶ **How:** by using information from previous gradient computations
 - ▶ For $\delta \in [0, 1[$, at each epoch, v is updated by $v \leftarrow \delta \cdot v - \eta \nabla_{\theta} \mathcal{C}(\theta)$



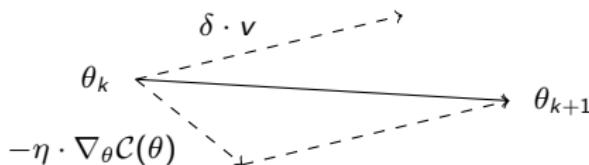
Momentum

- ▶ Principle: consider a *velocity* variable v
 - ▶ **Intuition:** v provides the direction and speed at which θ should move toward the optimum
 - ▶ **How:** by using information from previous gradient computations
 - ▶ For $\delta \in [0, 1[$, at each epoch, v is updated by $v \leftarrow \delta \cdot v - \eta \nabla_{\theta} \mathcal{C}(\theta)$



Momentum

- ▶ Principle: consider a *velocity* variable v
 - ▶ **Intuition:** v provides the direction and speed at which θ should move toward the optimum
 - ▶ **How:** by using information from previous gradient computations
 - ▶ For $\delta \in [0, 1[$, at each epoch, v is updated by $v \leftarrow \delta \cdot v - \eta \nabla_{\theta} \mathcal{C}(\theta)$



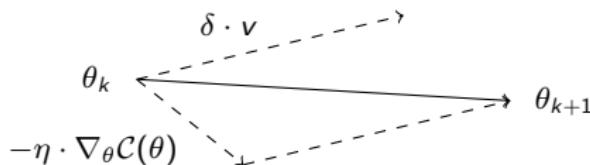
- ▶ Algorithm:
- Input:** Learning rate η , momentum parameter δ
Input: Initial parameters θ , initial velocity $v = 0$

```

1 for i ← 1 to number of training steps do
2   g ← computation of  $\nabla_{\theta} \mathcal{C}(\theta_k)$  //full gradient, mini-batch...;
3   v ←  $\delta \cdot v - \eta \cdot g$ ;
4    $\theta \leftarrow \theta + v$ 
5 end
  
```

Momentum

- ▶ Principle: consider a *velocity* variable v
 - ▶ **Intuition:** v provides the direction and speed at which θ should move toward the optimum
 - ▶ **How:** by using information from previous gradient computations
 - ▶ For $\delta \in [0, 1[$, at each epoch, v is updated by $v \leftarrow \delta \cdot v - \eta \nabla_{\theta} \mathcal{C}(\theta)$



- ▶ Algorithm:
 - Input:** Learning rate η , momentum parameter δ
 - Input:** Initial parameters θ , initial velocity $v = 0$

```

1 for i ← 1 to number of training steps do
2   g ← computation of  $\nabla_{\theta} \mathcal{C}(\theta_k)$  //full gradient, mini-batch...;
3   v ←  $\delta \cdot v - \eta \cdot g$ ;
4    $\theta \leftarrow \theta + v$ 
5 end
  
```

- ▶ Demo 7

Nesterov momentum

- ▶ Principle: use a *modified* velocity variable v

Nesterov momentum

- ▶ Principle: use a *modified* velocity variable v
 - ▶ Intuition: look ahead with current velocity before computing gradient

Nesterov momentum

- ▶ Principle: use a *modified* velocity variable v
 - ▶ Intuition: look ahead with current velocity before computing gradient
 - ▶ For $\delta \in [0, 1[$, at each epoch, v is updated by $v \leftarrow \delta \cdot v - \eta \nabla_{\theta} \mathcal{C}(\theta + \delta \cdot v)$

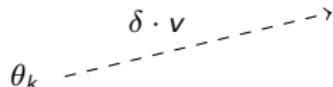
Nesterov momentum

- ▶ Principle: use a *modified* velocity variable v
 - ▶ Intuition: look ahead with current velocity before computing gradient
 - ▶ For $\delta \in [0, 1[$, at each epoch, v is updated by $v \leftarrow \delta \cdot v - \eta \nabla_{\theta} C(\theta + \delta \cdot v)$

$$\theta_k$$

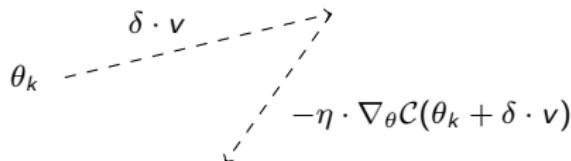
Nesterov momentum

- ▶ Principle: use a *modified* velocity variable v
 - ▶ Intuition: look ahead with current velocity before computing gradient
 - ▶ For $\delta \in [0, 1[$, at each epoch, v is updated by $v \leftarrow \delta \cdot v - \eta \nabla_{\theta} \mathcal{C}(\theta + \delta \cdot v)$



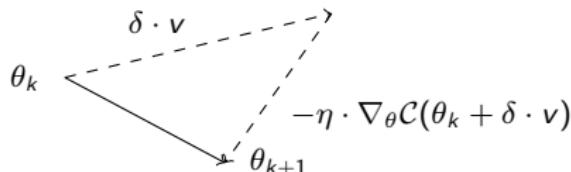
Nesterov momentum

- ▶ Principle: use a *modified* velocity variable v
 - ▶ Intuition: look ahead with current velocity before computing gradient
 - ▶ For $\delta \in [0, 1[$, at each epoch, v is updated by $v \leftarrow \delta \cdot v - \eta \nabla_{\theta} \mathcal{C}(\theta + \delta \cdot v)$



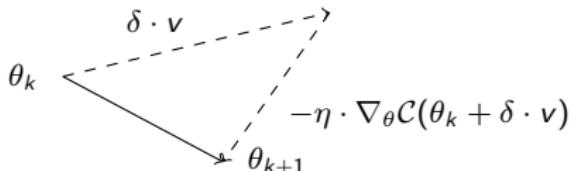
Nesterov momentum

- ▶ Principle: use a *modified* velocity variable v
 - ▶ Intuition: look ahead with current velocity before computing gradient
 - ▶ For $\delta \in [0, 1[$, at each epoch, v is updated by $v \leftarrow \delta \cdot v - \eta \nabla_{\theta} \mathcal{C}(\theta + \delta \cdot v)$



Nesterov momentum

- ▶ Principle: use a *modified* velocity variable v
 - ▶ Intuition: look ahead with current velocity before computing gradient
 - ▶ For $\delta \in [0, 1[$, at each epoch, v is updated by $v \leftarrow \delta \cdot v - \eta \nabla_{\theta} \mathcal{C}(\theta + \delta \cdot v)$



- ▶ Algorithm:
- Input:** Learning rate η , momentum parameter δ
Input: Initial parameters θ , initial velocity $v = 0$

```

1 for i ← 1 to number of training steps do
2   |    $\tilde{\theta} \leftarrow \theta + \delta \cdot v;$ 
3   |    $g \leftarrow \text{computation of } \nabla_{\theta} \mathcal{C}(\tilde{\theta}) \text{ //full gradient, mini-batch...};$ 
4   |    $v \leftarrow \delta \cdot v - \eta \cdot g;$ 
5   |    $\theta \leftarrow \theta + v;$ 
6 end

```

Implementing Nesterov Accelerated Gradient

- ▶ Update rules:

$$\begin{aligned}v &\leftarrow \delta \cdot v - \eta \nabla_{\theta} \mathcal{C}(\theta + \delta \cdot v) \\ \theta &\leftarrow \theta + v\end{aligned}$$

Implementing Nesterov Accelerated Gradient

- ▶ Update rules:

$$\begin{aligned}v &\leftarrow \delta \cdot v - \eta \nabla_{\theta} \mathcal{C}(\theta + \delta \cdot v) \\ \theta &\leftarrow \theta + v\end{aligned}$$

- ▶ The gradient computation is *forward looking*

Implementing Nesterov Accelerated Gradient

- ▶ Update rules:

$$\begin{aligned}v &\leftarrow \delta \cdot v - \eta \nabla_{\theta} \mathcal{C}(\theta + \delta \cdot v) \\ \theta &\leftarrow \theta + v\end{aligned}$$

- ▶ The gradient computation is *forward looking*
- ▶ This does not fit well in a generic implementation of forward and backpropagation algorithms

Implementing Nesterov Accelerated Gradient

- ▶ Update rules:

$$\begin{aligned}v &\leftarrow \delta \cdot v - \eta \nabla_{\theta} \mathcal{C}(\theta + \delta \cdot v) \\ \theta &\leftarrow \theta + v\end{aligned}$$

- ▶ The gradient computation is *forward looking*
- ▶ This does not fit well in a generic implementation of forward and backpropagation algorithms
- ▶ **A solution:** “simplified” Nesterov momentum (Bengio et al.: *Advances in optimizing recurrent networks*. 2012)

Simplified Nesterov momentum

- ▶ Perform updates on the lookahead variable $\Theta \stackrel{\text{def}}{=} \theta + \delta \cdot v$

Simplified Nesterov momentum

- ▶ Perform updates on the lookahead variable $\Theta \stackrel{\text{def}}{=} \theta + \delta \cdot v$
- ▶ At iteration $k + 1$ we have:

$$v_{k+1} = \delta \cdot v_k - \eta \cdot \nabla_\theta \mathcal{C}(\Theta_k)$$

Simplified Nesterov momentum

- ▶ Perform updates on the lookahead variable $\Theta \stackrel{\text{def}}{=} \theta + \delta \cdot v$
- ▶ At iteration $k + 1$ we have:

$$\begin{aligned}v_{k+1} &= \delta \cdot v_k - \eta \cdot \nabla_\theta \mathcal{C}(\Theta_k) \\ \Theta_{k+1} &= \theta_{k+1} + \delta \cdot v_{k+1}\end{aligned}$$

Simplified Nesterov momentum

- ▶ Perform updates on the lookahead variable $\Theta \stackrel{\text{def}}{=} \theta + \delta \cdot v$
- ▶ At iteration $k + 1$ we have:

$$\begin{aligned}v_{k+1} &= \delta \cdot v_k - \eta \cdot \nabla_\theta \mathcal{C}(\Theta_k) \\ \Theta_{k+1} &= \theta_{k+1} + \delta \cdot v_{k+1} = (\theta_k + v_{k+1}) + \delta \cdot v_{k+1}\end{aligned}$$

Simplified Nesterov momentum

- ▶ Perform updates on the lookahead variable $\Theta \stackrel{\text{def}}{=} \theta + \delta \cdot v$
- ▶ At iteration $k+1$ we have:

$$\begin{aligned}v_{k+1} &= \delta \cdot v_k - \eta \cdot \nabla_\theta \mathcal{C}(\Theta_k) \\ \Theta_{k+1} &= \theta_{k+1} + \delta \cdot v_{k+1} = (\theta_k + v_{k+1}) + \delta \cdot v_{k+1} \\ &= (\Theta_k - \delta \cdot v_k + v_{k+1}) + \delta \cdot v_{k+1}\end{aligned}$$

Simplified Nesterov momentum

- ▶ Perform updates on the lookahead variable $\Theta \stackrel{\text{def}}{=} \theta + \delta \cdot v$
- ▶ At iteration $k + 1$ we have:

$$\begin{aligned}v_{k+1} &= \delta \cdot v_k - \eta \cdot \nabla_\theta \mathcal{C}(\Theta_k) \\ \Theta_{k+1} &= \theta_{k+1} + \delta \cdot v_{k+1} = (\theta_k + v_{k+1}) + \delta \cdot v_{k+1} \\ &= (\Theta_k - \delta \cdot v_k + v_{k+1}) + \delta \cdot v_{k+1} \\ &= \Theta_k - \delta \cdot v_k + (1 + \delta) \cdot v_{k+1}\end{aligned}$$

Simplified Nesterov momentum

- ▶ Perform updates on the lookahead variable $\Theta \stackrel{\text{def}}{=} \theta + \delta \cdot v$
- ▶ At iteration $k + 1$ we have:

$$\begin{aligned}v_{k+1} &= \delta \cdot v_k - \eta \cdot \nabla_\theta \mathcal{C}(\Theta_k) \\ \Theta_{k+1} &= \theta_{k+1} + \delta \cdot v_{k+1} = (\theta_k + v_{k+1}) + \delta \cdot v_{k+1} \\ &= (\Theta_k - \delta \cdot v_k + v_{k+1}) + \delta \cdot v_{k+1} \\ &= \Theta_k - \delta \cdot v_k + (1 + \delta) \cdot v_{k+1} \\ &= \Theta_k - \delta \cdot v_k + (1 + \delta) \cdot (\delta \cdot v_k - \eta \cdot \nabla_\theta \mathcal{C}(\Theta_k))\end{aligned}$$

Simplified Nesterov momentum

- ▶ Perform updates on the lookahead variable $\Theta \stackrel{\text{def}}{=} \theta + \delta \cdot v$
- ▶ At iteration $k + 1$ we have:

$$\begin{aligned}v_{k+1} &= \delta \cdot v_k - \eta \cdot \nabla_{\theta} \mathcal{C}(\Theta_k) \\ \Theta_{k+1} &= \theta_{k+1} + \delta \cdot v_{k+1} = (\theta_k + v_{k+1}) + \delta \cdot v_{k+1} \\ &= (\Theta_k - \delta \cdot v_k + v_{k+1}) + \delta \cdot v_{k+1} \\ &= \Theta_k - \delta \cdot v_k + (1 + \delta) \cdot v_{k+1} \\ &= \Theta_k - \delta \cdot v_k + (1 + \delta) \cdot (\delta \cdot v_k - \eta \cdot \nabla_{\theta} \mathcal{C}(\Theta_k)) \\ &= \Theta_k + \delta^2 \cdot v_k - \eta \cdot (1 + \delta) \cdot \nabla_{\theta} \mathcal{C}(\Theta_k)\end{aligned}$$

Simplified Nesterov momentum

- ▶ Perform updates on the lookahead variable $\Theta \stackrel{\text{def}}{=} \theta + \delta \cdot v$
- ▶ At iteration $k + 1$ we have:

$$\begin{aligned}
 v_{k+1} &= \delta \cdot v_k - \eta \cdot \nabla_\theta \mathcal{C}(\Theta_k) \\
 \Theta_{k+1} &= \theta_{k+1} + \delta \cdot v_{k+1} = (\theta_k + v_{k+1}) + \delta \cdot v_{k+1} \\
 &= (\Theta_k - \delta \cdot v_k + v_{k+1}) + \delta \cdot v_{k+1} \\
 &= \Theta_k - \delta \cdot v_k + (1 + \delta) \cdot v_{k+1} \\
 &= \Theta_k - \delta \cdot v_k + (1 + \delta) \cdot (\delta \cdot v_k - \eta \cdot \nabla_\theta \mathcal{C}(\Theta_k)) \\
 &= \Theta_k + \delta^2 \cdot v_k - \eta \cdot (1 + \delta) \cdot \nabla_\theta \mathcal{C}(\Theta_k)
 \end{aligned}$$

- ▶ Algorithm:

Input: Learning rate η , momentum parameter δ
Input: Initial parameters θ , initial velocity $v = 0$

```

1  $\Theta \leftarrow \theta;$ 
2 for  $i \leftarrow 1$  to number of training steps do
3    $g \leftarrow$  computation of  $\nabla_\theta \mathcal{C}(\Theta)$  //full gradient, mini-batch...
4    $\Theta \leftarrow \Theta + \delta^2 \cdot v - \eta \cdot (1 + \delta) \cdot g;$ 
5    $v \leftarrow \delta \cdot v - \eta \cdot g;$ 
6 end

```

Simplified Nesterov momentum

- ▶ Perform updates on the lookahead variable $\Theta \stackrel{\text{def}}{=} \theta + \delta \cdot v$
- ▶ At iteration $k + 1$ we have:

$$\begin{aligned}
 v_{k+1} &= \delta \cdot v_k - \eta \cdot \nabla_\theta \mathcal{C}(\Theta_k) \\
 \Theta_{k+1} &= \theta_{k+1} + \delta \cdot v_{k+1} = (\theta_k + v_{k+1}) + \delta \cdot v_{k+1} \\
 &= (\Theta_k - \delta \cdot v_k + v_{k+1}) + \delta \cdot v_{k+1} \\
 &= \Theta_k - \delta \cdot v_k + (1 + \delta) \cdot v_{k+1} \\
 &= \Theta_k - \delta \cdot v_k + (1 + \delta) \cdot (\delta \cdot v_k - \eta \cdot \nabla_\theta \mathcal{C}(\Theta_k)) \\
 &= \Theta_k + \delta^2 \cdot v_k - \eta \cdot (1 + \delta) \cdot \nabla_\theta \mathcal{C}(\Theta_k)
 \end{aligned}$$

- ▶ Algorithm:

Input: Learning rate η , momentum parameter δ

Input: Initial parameters θ , initial velocity $v = 0$

```

1  $\Theta \leftarrow \theta;$ 
2 for  $i \leftarrow 1$  to number of training steps do
3    $g \leftarrow$  computation of  $\nabla_\theta \mathcal{C}(\Theta)$  //full gradient, mini-batch...
4    $\Theta \leftarrow \Theta + \delta^2 \cdot v - \eta \cdot (1 + \delta) \cdot g;$ 
5    $v \leftarrow \delta \cdot v - \eta \cdot g;$ 
6 end

```

- ▶ If the optimum is reached at step n ($v_n = 0$) then $\Theta_n = \theta_n$

Which optimization to choose?

- ▶ No clear answer

Which optimization to choose?

- ▶ No clear answer
- ▶ Momentum and Adam are quite popular

Which optimization to choose?

- ▶ No clear answer
- ▶ Momentum and Adam are quite popular
- ▶ Recommendation: choose an optimizer and practice tuning its hyperparameters

Which optimization to choose?

- ▶ No clear answer
- ▶ Momentum and Adam are quite popular
- ▶ Recommendation: choose an optimizer and practice tuning its hyperparameters
- ▶ Some reading material
 - ▶ Momentum: <https://distill.pub/2017/momentum/>
 - ▶ Adam: <https://arxiv.org/abs/1412.6980>